

**A Report
On
Assignment
CSE 465**

By

Name: Fahmida Akter Nourin

Id: 2031741642

Question 1 (A):

```
Python
import torch
import torch.nn as nn
from torchvision.transforms import ToTensor

class task_A(nn.Module):
    def __init__(self):
        super(task_A, self).__init__()
        self.conv1x1 = nn.Conv2d(256, 128, kernel_size=1)
        self.conv3x3 = nn.Conv2d(256, 192, kernel_size=3, padding=1)
        self.conv5x5 = nn.Conv2d(256, 96, kernel_size=5, padding=2)
        self.pool3x3 = nn.MaxPool2d(kernel_size=3, stride=1, padding=1)

    def forward(self, x):
        x1 = nn.ReLU()(self.conv1x1(x))
        x2 = nn.ReLU()(self.conv3x3(x))
        x3 = nn.ReLU()(self.conv5x5(x))
        pooled_x = self.pool3x3(x)
        out = torch.cat((x1, x2, x3, pooled_x), dim=1)
        return out

input_tensor = torch.randn(1, 256, 28, 28)
model = task_A()
output_tensor = model(input_tensor)

print(model)
print(output_tensor.shape)
```

Output:

```
Python
Output:
task_A(
  (conv1x1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1))
  (conv3x3): Conv2d(256, 192, kernel_size=(3, 3), stride=(1, 1), padding=(1,
  1))
  (conv5x5): Conv2d(256, 96, kernel_size=(5, 5), stride=(1, 1), padding=(2,
  2))
  (pool3x3): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1,
  ceil_mode=False)
)
torch.Size([1, 672, 28, 28])
```

Question 1 (B):

```
Python
class task_B(nn.Module):
    def __init__(self):
        super(task_B, self).__init__()
        self.branch1 = nn.Conv2d(256, 128, kernel_size = 1)
        self.branch2 = nn.Sequential(
            nn.Conv2d(256, 64, kernel_size = 1),
            nn.ReLU(),
            nn.Conv2d(64, 192, kernel_size = 3, padding = 1),
        )
        self.branch3 = nn.Sequential(
            nn.Conv2d(256, 64, kernel_size = 1),
            nn.ReLU(),
            nn.Conv2d(64, 96, kernel_size = 5, padding = 2),
        )
        self.branch4 = nn.Sequential(
            nn.MaxPool2d(kernel_size = 3, stride = 1, padding = 1),
            nn.Conv2d(256, 64, kernel_size = 1),
        )

    def forward(self, x):
        b1_out = self.branch1(x)
        b2_out = self.branch2(x)
        b3_out = self.branch3(x)
        b4_out = self.branch4(x)
        out = torch.cat((b1_out, b2_out, b3_out, b4_out), dim=1)
        return out

input_x = torch.randn(1, 256, 28, 28)
model = task_B()
output = model(input_x)
print(model)
print(output.shape)
```

Output:

```
Python
task_B(
  (branch1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1))
  (branch2): Sequential(
    (0): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1))
```

```

(1): ReLU()
(2): Conv2d(64, 192, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
)
(branch3): Sequential(
(0): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1))
(1): ReLU()
(2): Conv2d(64, 96, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
)
(branch4): Sequential(
(0): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1, ceil_mode=False)
(1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1))
)
)
torch.Size([1, 480, 28, 28])

```

Question 1 (C):

Python

```

import torch
import torch.nn as nn

class RNN(nn.Module):
    def __init__(self, input_size, hidden_size, num_layers=1, dropout=0.1):
        super().__init__()
        self.hidden_size = hidden_size
        self.num_layers = num_layers
        self.rnn = nn.RNN(input_size, hidden_size, num_layers, dropout=dropout)

    def forward(self, input, hidden):
        output, hidden = self.rnn(input, hidden)
        return output


class Softmax(nn.Module):
    def __init__(self, input_size):
        super().__init__()
        self.linear = nn.Linear(input_size, input_size)

    def forward(self, input):
        output = self.linear(input)

```

```

probs = torch.nn.functional.softmax(output, dim=1)
return probs

# Convert the input text to word embeddings
input_text = "The quick brown fox"
word_to_index = {word: idx for idx, word in enumerate(input_text.lower().split())}
input_sequence = torch.tensor([[word_to_index[word] for word in
input_text.lower().split()]] , dtype=torch.float)

input_size = 4 #len(input_text.split(" "))
hidden_layers_size = 4 #len(input_text.split(" "))
num_layers = 2
dropout = 0.2

rnn_module = RNN(input_size, hidden_layers_size, num_layers, dropout)
softmax_module = Softmax(hidden_layers_size)

# Initialize hidden state with correct shape
hidden_layers = torch.zeros(num_layers, hidden_layers_size, dtype=torch.float)

# Pass input through modules:
output = rnn_module(input_sequence, hidden_layers)
prob = softmax_module(output)

print(prob)

```

OutPut

Python

```

tensor([[0.3461, 0.2818, 0.1557, 0.2164]], grad_fn=<SoftmaxBackward0>)

```

Question 2

```
Python

import torch
import torch.nn as nn
import torchvision.models as models
from torchvision import transforms, datasets
from PIL import Image
import matplotlib.pyplot as plt
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

img_path = '/kaggle/input/vegetable-image-dataset/Vegetable
Images/train/Bean/0026.jpg'
img = Image.open(img_path).convert('RGB')

#transforming image
preprocess = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
])

img_tensor = preprocess(img)
img_tensor = torch.unsqueeze(img_tensor, 0).to(device)

transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
])

train_dataset =
datasets.ImageFolder(root='/kaggle/input/vegetable-image-dataset/Vegetable
Images/train', transform=transform)
train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=32,
shuffle=True)

num_classes = 15
```

```

# function of the feature map of the specified layer
def plot_feature_map(model, img_tensor, idx):
    feature_map_model = torch.nn.Sequential(*list(model.children())[:idx])
    feature_map_model.to(device)
    feature_map = feature_map_model(img_tensor)

    with torch.no_grad():
        feature_map = feature_map_model(img_tensor)

    # Visualize the feature map
    plt.imshow(feature_map[0, 0].cpu().detach().numpy(), cmap='viridis')
    plt.show()

def Simple_train(train_loader, model, criterion, optimizer):
    num_epochs = 3
    for epoch in range(num_epochs):
        for inputs, labels in train_loader:
            inputs, labels = inputs.to(device), labels.to(device)

            optimizer.zero_grad()
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

        print(f'Epoch {epoch + 1}/{num_epochs}, Loss: {loss.item()}')
    return model

```

Python

```

vgg16_model = models.vgg16(pretrained=True)
vgg16_model.classifier[-1] = nn.Linear(4096, num_classes)

# Train the model VGG16
vgg16_model.to(device)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(vgg16_model.parameters(), lr=0.001, momentum=0.9)

```

```
vgg16_model.train()
vgg16_model = Simple_train(train_loader, vgg16_model, criterion, optimizer)
```

Python

```
resnet101_model = models.resnet101(pretrained=True)
resnet101_model.fc = nn.Linear(resnet101_model.fc.in_features, num_classes)

# Train the model resnet10
resnet101_model.to(device)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(resnet101_model.parameters(), lr=0.001, momentum=0.9)
resnet101_model.train()
resnet101_model = Simple_train(train_loader, resnet101_model, criterion, optimizer)
```

Python

```
mobilenetv3_model = models.mobilenet_v3_small(pretrained=True)
mobilenetv3_model.classifier[-1] =
nn.Linear(mobilenetv3_model.classifier[-1].in_features, num_classes)

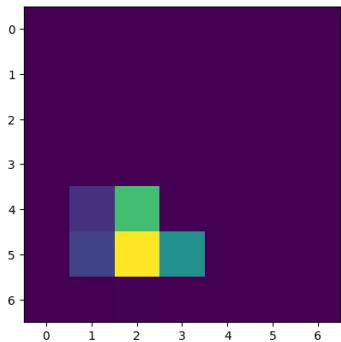
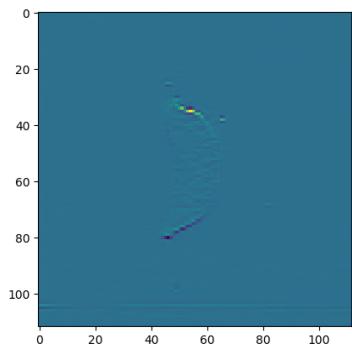
# Train the model mobilen
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(mobilenetv3_model.parameters(), lr=0.001,
momentum=0.9)

mobilenetv3_model.to(device)
mobilenetv3_model.train()
mobilenetv3_model =
Simple_train(train_loader, mobilenetv3_model, criterion, optimizer)
```

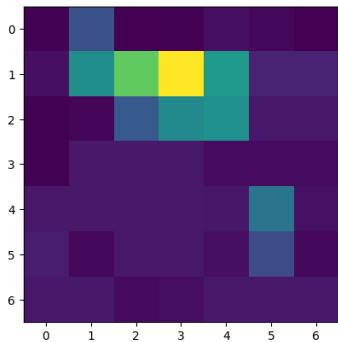
(a)

Python

```
print("feature map of the first level features using Resnet-101")
plot_feature_map(resnet101_model,img_tensor,1)
print("feature map of the first level features using VGG16")
plot_feature_map(vgg16_model,img_tensor,1)
print("feature map of the first level features using MobilenetV3")
plot_feature_map(mobilenetv3_model,img_tensor,1)
```



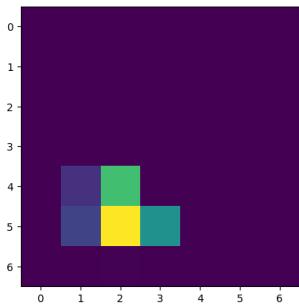
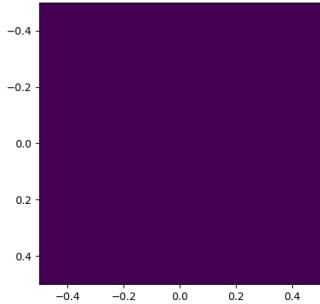
Actual Image



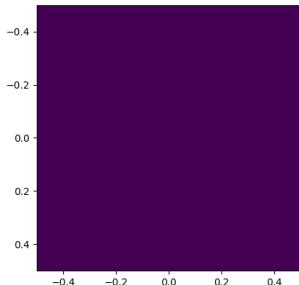
(b)

Python

```
print("feature map of the final level features using Resnet-101")
plot_feature_map(resnet101_model,img_tensor,-1)
print("feature map of the first level features using VGG16")
plot_feature_map(vgg16_model,img_tensor,-1)
print("feature map of the final level features using MobilenetV3")
plot_feature_map(mobilenetv3_model,img_tensor,-1)
```



Actual Image



2(c)

```
Python

import numpy as np
from sklearn.neighbors import NearestNeighbors
from pathlib import Path
from PIL import Image
from torchvision import transforms, models
import torch
from tqdm.notebook import tqdm
from PIL import Image
import matplotlib.pyplot as plt
import numpy as np

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

def plot_clg(image_paths):
    images = [Image.open(path) for path in image_paths]
    fig, axs = plt.subplots(1, 5, figsize=(10, 3))
    for i in range(len(images)):
        img_array = np.array(images[i])
        axs[i].imshow(img_array)
        axs[i].axis('off')
    plt.tight_layout()
    plt.show()

def extract_features(image_path, model):
    model.eval()
    preprocess = transforms.Compose([
        transforms.Resize((224, 224)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
    ])

    image = Image.open(image_path).convert('RGB')
    input_image = preprocess(image)
    input_batch = torch.unsqueeze(input_image, 0).to(device)

    with torch.no_grad():
        output = model(input_batch)
```

```
feature_vector = output.squeeze().cpu().numpy()

return feature_vector

vgg16_model = models.vgg16(pretrained=True).to(device)

test_image_dir = Path('/kaggle/input/vegetable-image-dataset/Vegetable
Images/test')
validation_image_dir = Path('/kaggle/input/vegetable-image-dataset/Vegetable
Images/validation')

print("Finding the lowest number image from each category in the test
set.....")
lowest_number_images = {}
for category_dir in tqdm(test_image_dir.iterdir()):
    if category_dir.is_dir():
        images_in_category = list(category_dir.iterdir())
        lowest_number_image = min(images_in_category, key=lambda x: int(x.stem))
        lowest_number_images[category_dir.name] = lowest_number_image

print("Extracting features from the lowest number
images.....")
query_features = {}
for category, image_path in tqdm(lowest_number_images.items()):
    features = extract_features(image_path, vgg16_model)
    query_features[category] = features

validation_features = []
validation_image_paths = []

print("Extracting features from all images in the validation
set.....")
for image_path in tqdm(validation_image_dir.glob('**/*.jpg')):
    features = extract_features(image_path, vgg16_model)
    validation_features.append(features)
    validation_image_paths.append(image_path)
```

```
query_features_array = np.array(list(query_features.values()))
validation_features_array = np.array(validation_features)

#find 5 nearest neighbors for each query
k_neighbors = 5
neigh = NearestNeighbors(n_neighbors=k_neighbors)
neigh.fit(validation_features_array)

nearest_neighbors_indices = []
for query_feature in query_features_array:
    _, indices = neigh.kneighbors([query_feature])
    nearest_neighbors_indices.append(indices.flatten())

for category, indices in zip(lowest_number_images.keys(),
nearest_neighbors_indices):
    print(f"Category: {category}")
    print("Nearest Neighbors:")
    temp=[]
    for index in indices:
        #print(validation_image_paths[index])
        temp.append(validation_image_paths[index])
    plot_clg(temp)
    print("\n")
```

Category: Broccoli

Nearest Neighbors:



Centered Image

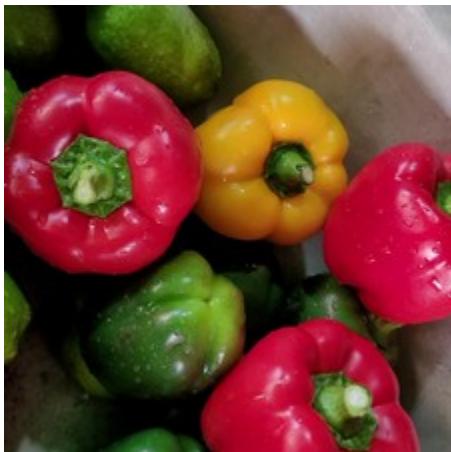


Category: Capsicum

Nearest Neighbors:



Centered Image



Category: Bottle_Gourd

Nearest Neighbors:



Centered Image



Category: Radish

Nearest Neighbors:



Centered Image



Category: Tomato

Nearest Neighbors:



Centered Image



Category: Brinjal

Nearest Neighbors:



Category: Pumpkin

Nearest Neighbors:

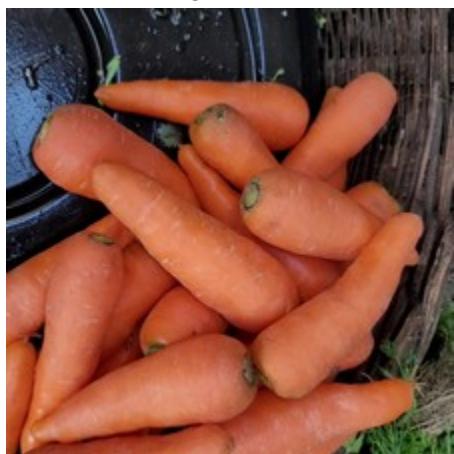


Category: Carrot

Nearest Neighbors:



Centered Image



Category: Papaya

Nearest Neighbors:



Category: Cabbage

Nearest Neighbors:



Centered Image



Category: Bitter_Gourd

Nearest Neighbors:



Actual Image



Category: Cauliflower

Nearest Neighbors:



Centered Image



Category: Bean

Nearest Neighbors:



Centered Image



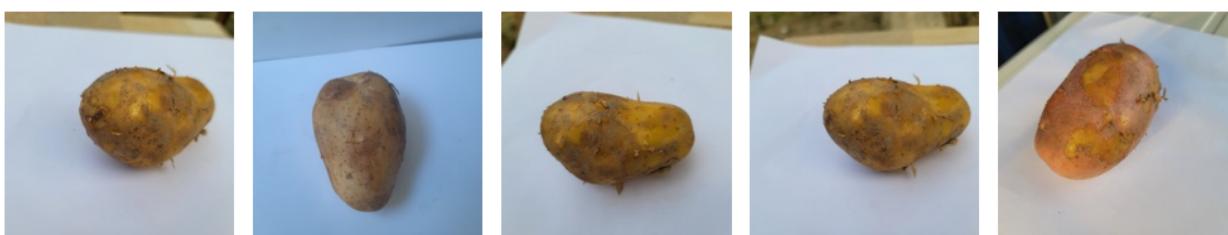
Category: Cucumber

Nearest Neighbors:



Category: Potato

Nearest Neighbors:



Centered Image



Question 3

Python

```
import torch
import torchvision
import torchvision.transforms as transforms
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from sklearn.metrics import confusion_matrix, accuracy_score
import time

# Defining data transformation
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
])

# Loading and validation datasets
train_dataset =
    torchvision.datasets.ImageFolder(root='/kaggle/input/vegetable-image-dataset/Vegetable_Images/test', transform=transform)
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True,
    num_workers=4)

val_dataset =
    torchvision.datasets.ImageFolder(root='/kaggle/input/vegetable-image-dataset/Vegetable_Images/validation', transform=transform)
```

```
val_loader = DataLoader(val_dataset, batch_size=32, shuffle=False, num_workers=4)

num_classes = 15

# Define VGG16 model
vgg_model = torchvision.models.vgg16(pretrained=True)
# Modifying the last layer
vgg_model.classifier[6] = nn.Linear(4096, num_classes)

# Define ResNet-101 model
resnet_model = torchvision.models.resnet101(pretrained=True)
#Modifying the last layer
resnet_model.fc = nn.Linear(2048, num_classes)
```

Python

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
vgg_model = vgg_model.to(device)
resnet_model = resnet_model.to(device)
```

Python

```
# Defining optimizer and loss function
criterion = nn.CrossEntropyLoss()
optimizer_vgg = optim.Adam(vgg_model.parameters(), lr=0.001)
optimizer_resnet = optim.SGD(resnet_model.parameters(), lr=0.01, momentum=0.9)
```

Python

```
# Training process
def train(model, train_loader, optimizer, criterion, device):
    model.train()
    running_loss = 0.0
```

```

correct_predictions = 0
total_samples = 0

for inputs, labels in train_loader:
    inputs, labels = inputs.to(device), labels.to(device)

    optimizer.zero_grad()

    outputs = model(inputs)
    loss = criterion(outputs, labels)
    loss.backward()
    optimizer.step()

    running_loss += loss.item()

    _, preds = torch.max(outputs, 1)
    correct_predictions += torch.sum(preds == labels.data)
    total_samples += labels.size(0)

accuracy = correct_predictions.double() / total_samples

return running_loss / len(train_loader), accuracy

```

```

# Validation process
def validate(model, val_loader, criterion, device):
    model.eval()
    all_preds = []
    all_labels = []

    with torch.no_grad():
        running_loss = 0.0
        correct_predictions = 0
        total_samples = 0

        for inputs, labels in val_loader:
            inputs, labels = inputs.to(device), labels.to(device)

            outputs = model(inputs)
            loss = criterion(outputs, labels)
            running_loss += loss.item()

            _, preds = torch.max(outputs, 1)

```

```

    correct_predictions += torch.sum(preds == labels.data)
    total_samples += labels.size(0)

    all_preds.extend(preds.cpu().numpy())
    all_labels.extend(labels.cpu().numpy())

accuracy = accuracy_score(all_labels, all_preds)
confusion_mat = confusion_matrix(all_labels, all_preds)

return running_loss / len(val_loader), accuracy, confusion_mat

```

Python

```

# Training loop
num_epochs = 5
start_time = time.time()

train_losses_vgg = []
val_losses_vgg = []
train_acc_vgg = []
val_acc_vgg = []

train_losses_resnet = []
val_losses_resnet = []
train_acc_resnet = []
val_acc_resnet = []

for epoch in range(num_epochs):

    train_loss_vgg_score, train_accuracy_vgg_score = train(vgg_model, train_loader,
optimizer_vgg, criterion, device)
    val_loss_vgg_score, val_acc_vgg_score, _ = validate(vgg_model, val_loader,
criterion, device)

    train_loss_resnet_score, train_accuracy_resnet_score = train(resnet_model,
train_loader, optimizer_resnet, criterion, device)
    val_loss_resnet_score, val_acc_resnet_score, _ = validate(resnet_model,
val_loader, criterion, device)

    train_losses_vgg.append(train_loss_vgg_score)

```

```

val_losses_vgg.append(val_loss_vgg_score)
train_acc_vgg.append(train_accuracy_vgg_score.item())
val_acc_vgg.append(val_acc_vgg_score.item())

train_losses_resnet.append(train_loss_resnet_score)
val_losses_resnet.append(val_loss_resnet_score)
train_acc_resnet.append(train_accuracy_resnet_score.item())
val_acc_resnet.append(val_acc_resnet_score)

print('-----')
-----')
print(f"Epoch {epoch+1}/{num_epochs} - VGG: Train Loss: {train_loss_vgg_score:.4f}, Train Acc: {train_accuracy_vgg_score:.4f}..... Val Loss: {val_loss_vgg_score:.4f}, Val Acc: {val_acc_vgg_score:.4f}")
print(f"Epoch {epoch+1}/{num_epochs} - ResNet: Train Loss: {train_loss_resnet_score:.4f}, Train Acc: {train_accuracy_resnet_score:.4f}..... Val Loss: {val_loss_resnet_score:.4f}, Val Acc: {val_acc_resnet_score:.4f}")

print('-----')
-----')
end_time = time.time()

```

output:

Python

```

-----
-----
-----
Epoch 1/5 - VGG: Train Loss: 2.6956, Train Acc: 0.1087..... Val Loss: 2.4158, Val Acc: 0.1653
Epoch 1/5 - ResNet: Train Loss: 0.4114, Train Acc: 0.8817..... Val Loss: 0.1997, Val Acc: 0.9527
-----
```

```
Epoch 2/5 - VGG: Train Loss: 2.1672, Train Acc: 0.2573 ..... Val Loss:  
1.9613, Val Acc: 0.3370  
Epoch 2/5 - ResNet: Train Loss: 0.1016, Train Acc: 0.9717 ..... Val Loss:  
0.0338, Val Acc: 0.9893
```

```
Epoch 3/5 - VGG: Train Loss: 1.7870, Train Acc: 0.3797 ..... Val Loss: 1.5250, Val Acc: 0.4777  
Epoch 3/5 - ResNet: Train Loss: 0.0503, Train Acc: 0.9860 ..... Val Loss: 0.0587, Val Acc: 0.9823
```

```
Epoch 4/5 - VGG: Train Loss: 1.4761, Train Acc: 0.4870..... Val Loss: 1.6112, Val Acc: 0.4380  
Epoch 4/5 - ResNet: Train Loss: 0.0166, Train Acc: 0.9957..... Val Loss: 0.0635, Val Acc: 0.9823
```

```
Epoch 5/5 - VGG: Train Loss: 1.4045, Train Acc: 0.5223 ..... Val Loss: 1.3185, Val Acc: 0.5647  
Epoch 5/5 - ResNet: Train Loss: 0.0101, Train Acc: 0.9973 ..... Val Loss: 0.0163, Val Acc: 0.9950
```

```
Python
# Calculate training time
training_time = end_time - start_time
print(f"Training Time: {training_time:.2f} seconds")
```

output:

```
Python
Training Time: 337.79 seconds
```

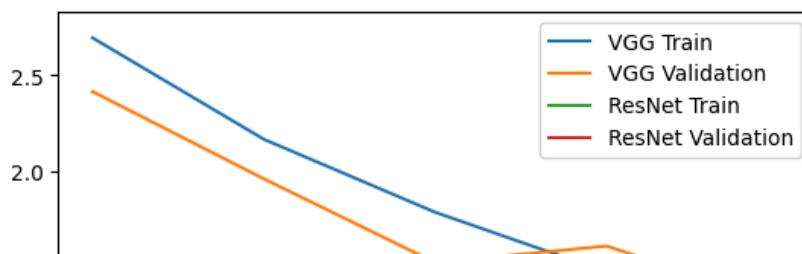
```
Python
# Print GPU memory footprint
print(f"Memory Usage: {torch.cuda.memory_allocated(device) / 1e9} GB")
```

ouput:

```
Python
Memory Usage: 2.69265664 GB
```

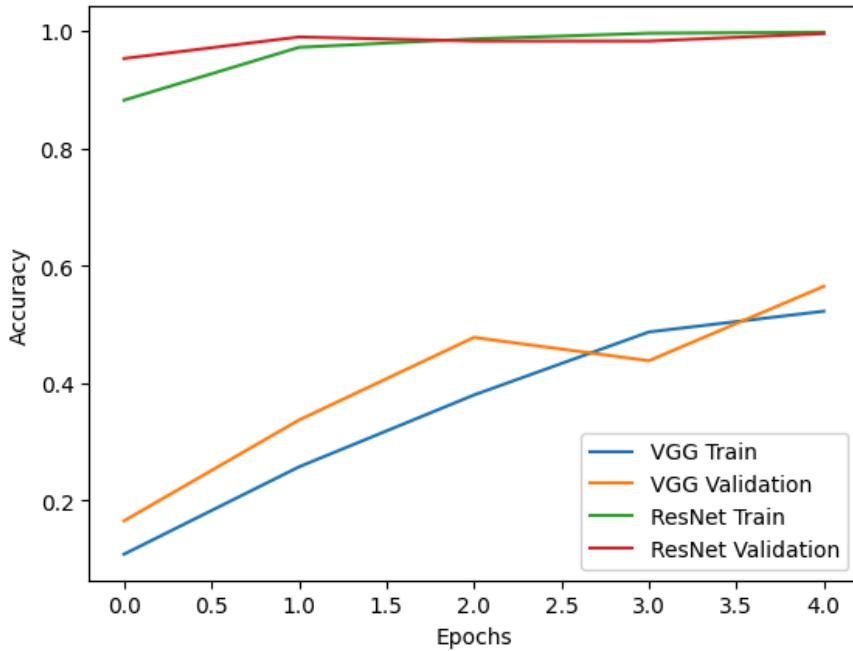
```
Python
# Plot training and validation loss curves
import matplotlib.pyplot as plt

plt.plot(train_losses_vgg, label='VGG Train')
plt.plot(val_losses_vgg, label='VGG Validation')
plt.plot(train_losses_resnet, label='ResNet Train')
plt.plot(val_losses_resnet, label='ResNet Validation')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



Python

```
import matplotlib.pyplot as plt
# Plot training and validation accuracy curves
plt.plot(train_acc_vgg, label='VGG Train')
plt.plot(val_acc_vgg, label='VGG Validation')
plt.plot(train_acc_resnet, label='ResNet Train')
plt.plot(val_acc_resnet, label='ResNet Validation')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



Python

```
# Testing process
def test(model, test_loader, criterion, device):
    model.eval()
    all_preds = []
    all_labels = []

    with torch.no_grad():
```

```

running_loss = 0.0
correct_predictions = 0
total_samples = 0

for inputs, labels in test_loader:
    inputs, labels = inputs.to(device), labels.to(device)

    outputs = model(inputs)
    loss = criterion(outputs, labels)
    running_loss += loss.item()

    _, preds = torch.max(outputs, 1)

    correct_predictions += torch.sum(preds == labels.data)
    total_samples += labels.size(0)

    all_preds.extend(preds.cpu().numpy())
    all_labels.extend(labels.cpu().numpy())

accuracy = accuracy_score(all_labels, all_preds)
confusion_mat = confusion_matrix(all_labels, all_preds)

return running_loss / len(test_loader), accuracy, confusion_mat

# Testing the models
test_dataset =
torchvision.datasets.ImageFolder(root='/kaggle/input/vegetable-image-dataset/Vegetable_Images/test', transform=transform)
test_loader = DataLoader(val_dataset, batch_size=32, shuffle=False, num_workers=4)

test_loss_vgg, test_acc_vgg, confusion_mat_vgg = test(vgg_model, test_loader,
criterion, device)
test_loss_resnet, test_acc_resnet, confusion_mat_resnet = test(resnet_model,
test_loader, criterion, device)

# Plot confusion matrix for VGG
print("VGG Test Results:")
print(f"Test Loss: {test_loss_vgg:.4f}, Test Accuracy: {test_acc_vgg:.4f}")
print("Confusion Matrix:")
plt.figure(figsize=(10, 8))
sns.heatmap(confusion_mat_vgg, annot=True, fmt="d", cmap="Blues",
xticklabels=range(num_classes), yticklabels=range(num_classes))
plt.xlabel("Predicted")
plt.ylabel("Actual")

```

```

plt.title("Confusion Matrix - VGG")
plt.show()

# Plot confusion matrix for Resnet
print("\nResNet Test Results:")
print(f"Test Loss: {test_loss_resnet:.4f}, Test Accuracy: {test_acc_resnet:.4f}")
print("Confusion Matrix:")
print(confusion_mat_resnet)

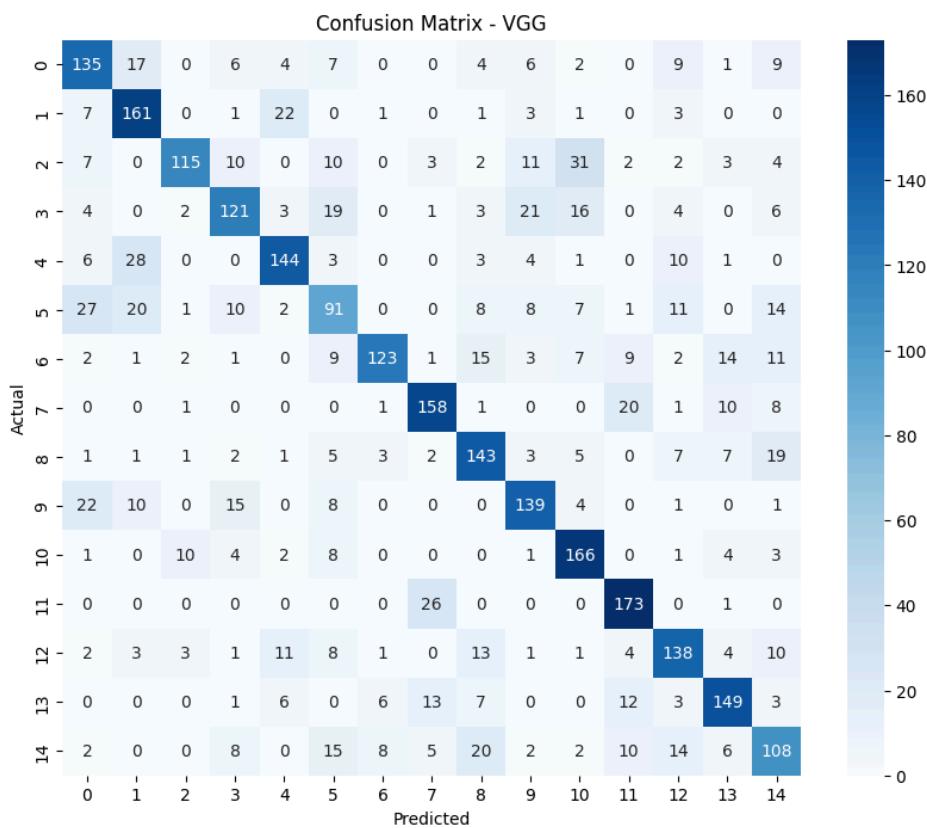
# Plot confusion matrix for ResNet
plt.figure(figsize=(10, 8))
sns.heatmap(confusion_mat_resnet, annot=True, fmt="d", cmap="Blues",
            xticklabels=range(num_classes), yticklabels=range(num_classes))
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix - ResNet")
plt.show()

```

VGG Test Results:

Test Loss: 0.9910, Test Accuracy: 0.6880

Confusion Matrix:



ResNet Test Results:

Test Loss: 0.0174, Test Accuracy: 0.9933

Confusion Matrix:

